

INDIA

SIBI CHAKKARAVARTHY SETHURAMAN · JUNE 2026

AI is not just ChatGPT on your phone

A plain-language briefing for government staff — exam leaks, portal disasters, and why that chat app on your desk is a security problem.



Exam leaks



Open buckets



ChatGPT habit



Deepfakes



Global threats



What to do

INDIA

What is AI, really?

Not magic. Not a person. A pattern-matching engine that guesses the next word — very convincingly.

Everything you type can leave your building

Consumer AI stores conversations on vendor servers. 'Delete chat' ≠ gone from everywhere.

It will answer even when it doesn't know

Hallucination isn't a bug — the system is built to sound fluent, not to admit ignorance.

It works for the company that runs it

Not for your department, not for the citizen whose Aadhaar you pasted in.

Rule of thumb: if you wouldn't fax it to a stranger in California, don't paste it into ChatGPT.



Data leaves India

US/EU servers



Guesses fluently

even when wrong



Owes you nothing

no duty of care

INDIA

We start with India because the wounds are fresh

NTA, CBSE, state commissions — and your own office laptop.

Crores of students bet their futures on portals that crash on result day

Paper leaks still travel on Telegram — human couriers, not hackers in hoodies

Finance Ministry told staff in Jan 2025: stop using ChatGPT on office machines

The leaks today need people. Tomorrow they may not.

NEET

CANCELLED '26

CBSE

53 BUCKETS

FinMin

AI ADVISORY

Your voters' kids sit these exams. Your portals hold their data.

INDIA

NEET 2024 — the leak the Supreme Court acknowledged

May 5: paper out before the exam ended

CBI traced it to Oasis School, Hazaribagh — principal + centre superintendent.

~155 students got direct benefit

Networks in Patna and Hazaribagh; coaching centres in the loop.

Too many perfect scores

Unusual count of 720/720; grace marks fight dragged for weeks.

Court said leak was real — but refused full re-test

Results stood. Students who didn't cheat still paid the trust tax.

Insider + WhatsApp + Telegram. Same playbook for decades.

NEET 2024

Hazaribagh Oasis School

155 direct beneficiaries

SC: leak undisputed

INDIA

NEET 2026 — entire exam cancelled. First time ever.

Fresh leak allegations while 2024 cases still in court

Students protested nationwide — years of prep nullified overnight

Strongroom security, logistics, coaching nexus: unchanged weaknesses

Public Examinations Act 2024 is law. Enforcement is still catching up.

A phone photo → OCR → Telegram bot → 50,000 subscribers. That pipeline can run without a human courier.

NEET 2026

Full exam cancelled

First time in history

Same playbook, worse outcome

INDIA

CUET, UGC-NET, SSC — same disease, different headlines

CUET: wrong papers handed out, registration servers dying mid-upload
Kanpur centre got the wrong set. Thousands re-tested.

UGC-NET June 2024: 9+ lakh candidates sent home — exam cancelled next day

Plot twist: CBI said one 'leak' was a teenager's fake screenshot
Hoaxes and real leaks now look identical. Verification before panic matters.

SSC CGL 2025: candidate held the mouse, answers filled themselves
Dhanbad centre — exam server linked to remote operator.

JEE Main 2025: result day = 500 errors for 13 lakh students

2024–26 INCIDENTS

NEET

Cancelled '26

UGC-NET

9L+ cancelled

CUET

Wrong papers

SSC CGL

Remote hijack

JEE

500 on result day

Railway

Rs 1.17Cr racket

INDIA

CBSE 2026 — master password in the frontend

Teen ethical hackers Nisarga Adhikary & Tirth Parmar opened DevTools and found the keys to the kingdom.

`const MASTER_PASSWORD = '...' — sitting in JavaScript anyone can read`

SQL injection paths, no IP restrictions, evaluator logins exposed

Public S3 buckets: answer sheets, question papers, scanned booklets — no login required

CBSE said 'test portal' at first. Then admitted gaps. IIT teams called in.

Vendor reportedly ran answer sheets through Google Gemini

Third-party AI on children's exam data. Think about that.

Vibecoded: ship fast, audit never, pray.

>_ cbse-osm-portal.js

```
const CONFIG = {  
  masterPassword: "CBSE@2026!",  
  skipIpCheck: true,  
};
```

 Hardcoded in frontend JS. Found in production.

INDIA

It's not just exams — your portals leak too

Income Tax e-Filing: change PAN in URL → see someone else's Aadhaar + bank account

135 million registered users. Classic missing auth check.

S3WaaS govt cloud: years of open buckets with vaccination + passport data

Maharashtra 2025: Police, MPSC, pension logins on the darknet

Same failure mode as CBSE — misconfiguration, not genius hackers

A .gov.in domain does not mean someone audited the code.

IDOR IN E-FILING

```
GET /api/taxpayer?pan=
AAAAA0000A → your data
BBBBB1111B → someone else's
```

☰ Same class of bug as CBSE S3 buckets

💡 135M users. One missing check.

INDIA

The ChatGPT habit in government offices

FinMin circular, Jan 2025: no ChatGPT / DeepSeek on official devices

Rajya Sabha, Mar 2025: no blanket ban — but confidentiality rules haven't gone anywhere

Samsung engineers pasted chip secrets into ChatGPT three times in 20 days
If Fortune 500 engineers do it, your clerk will too.

'I'm just cleaning up the wording' = uploading sovereign data to a US cloud

Never paste: Aadhaar, citizen complaints, draft circulars, meeting minutes with names, file numbers.



Clipboard ≠ private

FinMin advisory · Jan 2025

- X Aadhaar numbers
- X Draft circulars
- X Citizen complaints
- X File notings with names

INDIA

Deepfakes are already stealing money here

Fake PM / FM videos selling 'government-backed' crypto schemes

Go Invest, Cryptify — Times of India layout cloned.

Pune army veteran: Rs 1 crore gone to deepfake ministers + fake trading app

3 seconds of your kid's voice on Instagram → convincing 'kidnap' WhatsApp call

Delhi: Rs 50,000 lost. Kerala retiree: Rs 40,000.

Retired govt staff are the bullseye — they trust voice and uniform

1930 · cybercrime.gov.in — print it. Pin it. Repeat it.



Face + voice = trust

Cloned ministers. Cloned children. Real UPI losses.

Helpline 1930

INDIA

One small action → damage at scale

Paste a draft circular → vendor breach → text in a future model

One camera flash in a strongroom → Telegram → lakhs of futures altered

One password in source code → every answer sheet downloadable

One deepfake clip → thousands of UPI transfers before fact-checkers wake up

Today: humans at every step. Tomorrow: scripts connecting each step. The blast radius stays the same.

RIPPLE CHAIN



One paste
amplifies →



Portal gap
amplifies →



Exam leak
amplifies →



Deepfake

AI can script each step. Humans optional.

REASSURE

Pause — real help already exists

You don't need to become a cyber expert overnight. Government built channels for exactly this.

Step 1: Stop. Don't pay. Don't share OTPs.

Scammers rush you so you can't think. Taking 60 seconds is allowed.

Step 2: Call 1930 or report at cybercrime.gov.in

National helpline — banks may freeze fraudulent transfers if you act quickly.

Step 3: Walk to your IT desk or cyber cell the next morning

You are not bothering anyone. This is their job.

Step 4: Tell two colleagues what happened — they may be next

Sharing the story prevents panic, not embarrassment.

One reported scam can save a dozen colleagues from the same trap. You did the right thing.

You are not alone in this.

CALM STEPS

1 Stop — don't pay

2 Call 1930

3 Report online

4 Tell IT tomorrow

You are not bothering anyone. This is what these channels exist for.

REASSURE

Fake 'cyber police' vs the real thing

Scammers wear the uniform in a video. Real officers follow different rules.

Real: you report → you get a reference number → follow-up through official channels

Fake: urgent WhatsApp call, demands immediate UPI 'verification deposit'

Real: never asks for your ATM PIN, OTP, or remote screen-sharing app

Fake: threatens arrest in one hour unless you pay now

Fear is the product. Slow down.

When in doubt: hang up, call 1930, verify on cybercrime.gov.in

Trust the number you dial — not the number that dialled you.

POLICING & CYBERCRIME



cybercrime.gov.in

National reporting portal



Helpline 1930

Report fraud fast



Police logins leaked

Maharashtra Apr 2025



Fake cyber inspectors

WhatsApp impersonation

REASSURE

AI is already guarding your UPI payments

NPCI and RBI now run AI/ML models that score every suspicious transaction in real time — and have declined fraud before it left the account.

UPI fraud jumped 85% in FY 2023–24 — Rs 1,087 crore lost

NPCI's answer: stop reacting after the theft, start predicting it.

AI assigns a risk score to mule accounts based on money-flow history

When a fresh account suddenly collects from many people, the model flags it.

Banks get an alert → they call you → 'why this transfer?' before it clears

NPCI says it has already declined fraudulent transactions this way.

RBI's upcoming DPIP will score payments using telecom + cybercrime data

Mule accounts and synthetic IDs spotted before the money moves.

This is AI on the citizen's side — quietly, millions of times a day.

INDIA · UPI DEFENCE

Real-time

AI risk score on suspicious UPI transfers



Flags mule accounts by money-flow pattern



Bank calls you before the transfer clears



Fraud transactions already declined

NPCI / RBI (2025) — The Hindu BusinessLine, Medianama, ET BFSI

INDIA

Your 'anonymous' posts aren't anonymous

From the deanonymizer project — fusing small clues across platforms.

Employer, city, schedule, writing style — stitched together by an LLM

'Just venting about work' can name your department and your tehsil

LinkedIn + Facebook + that forum you forgot about = enough for targeted phishing

Run live exposure audit →

SIGNAL FUSION

"GST camp at collectorate"

+ "TT Nagar branch"

+ posting timestamps

= **your department, your city**

REASSURE

AI caught a real attack before the hackers could strike

Google's 'Big Sleep' AI agent found a critical SQLite flaw that only criminals knew about — and got it patched first.

Nov 2024: Big Sleep found the first-ever zero-day discovered by an AI agent

An exploitable memory bug in SQLite — software running inside billions of devices.

July 2025: it found CVE-2025-6965 — a flaw attackers were about to use

Google says this is the first time AI directly foiled a real in-the-wild exploit.

Same idea protects the boring software your portals quietly depend on

SQLite sits under countless government and banking apps.

The lesson: AI doesn't only help attackers — it gives defenders a head start

Found and fixed before users were ever at risk.

GOOGLE · BIG SLEEP

1st

zero-day ever found by an AI agent



Scans real code for unknown flaws



Caught CVE-2025-6965 before attackers struck



Patched before any user was at risk

Google Project Zero & DeepMind (2024–2025)

REASSURE

Machines that find and fix bugs — proven at DEF CON

DARPA's AI Cyber Challenge (Aug 2025) put autonomous AI systems against real code. They delivered.

AI systems scanned 54 million lines of code for vulnerabilities

Found 54 planted bugs and 18 real, unknown zero-days — then auto-wrote patches

Average fix time: about 45 minutes per vulnerability.

Winners (Team Atlanta, Trail of Bits, Theori) shared \$8.5M — and open-sourced everything

DARPA will point these tools at real critical-infrastructure code, including health care.

This is the same kind of audit India's exam and citizen portals desperately need

Imagine this run against a portal before launch — instead of teen hackers finding the master password after.

🛡️ DARPA · AIXCC · DEF CON 2025

18

real zero-days found by autonomous AI



54M lines of code scanned



Auto-patched in ~45 min each



Winning systems open-sourced

DARPA AIXCC results (Aug 2025) — arpa.mil, CyberScoop

REASSURE

When attackers used AI, defenders' AI caught them

Nov 2025: Anthropic detected and shut down the first AI-run espionage campaign — in about 10 days.

A state-linked group tricked an AI agent into doing 80–90% of an attack

Reconnaissance and exploitation at thousands of requests per second — impossible for humans.

Anthropic's threat-intelligence team spotted the abuse and disrupted it

Accounts banned, ~30 targeted organisations warned, defenses upgraded.

The AI even 'hallucinated' fake credentials — slowing the attackers down

Takeaway: monitoring + AI defenders can match AI-speed attacks

The defenders are using the same tools — and they're watching.

ANTHROPIC · THREAT INTEL

~10 days

to detect & shut an AI-run attack



Spotted AI doing 80–90% of an attack



~30 targeted orgs warned



Accounts banned, defences upgraded

Anthropic report (Nov 2025) — MITRE ATT&CK C0062

REASSURE

So what can AI safely do at your desk?

Think of it as a fast junior clerk who never sleeps — but who shows every draft to you.

Summarise long public circulars into plain language for citizens

Draft first versions of replies and translations — you edit, approve, sign

Tamil, Telugu, Marathi, Bengali notices with human review.

Answer routine staff questions at night via a guarded internal bot

This workshop's /demo — try it after the session.

Check your own public footprint defensively

/demo?tab=audit — see what strangers could infer about you.

Positive AI = approved tool + human sign-off + audit trail. You stay the decision-maker.

POSITIVE USES



Guarded staff chatbot
Refusals + Indian context



Translation & plain language
Human review before send



Fraud pattern detection
Supports investigators



Vendor audit before go-live
Not vibecoded CBSE

REASSURE

The machine proposes — you dispose

That is the whole governance model. AI doesn't change who is accountable.

AI drafts → you read → you edit → you approve → you send

AI flags a risky transaction → a human officer makes the call

Exactly how NPCI's fraud alerts work — the bank still phones you.

If output feels wrong, discard it — no penalty, no machine override

No honest vendor sells 'replace the officer' — only 'assist the officer'

Awareness ≠ anxiety. You now know both the risks and the tools fighting them.

CALM STEPS

1 Stop — don't pay

2 Call 1930

3 Report online

4 Tell IT tomorrow

You are not bothering anyone. This is what these channels exist for.

GLOBAL

Zoom out — the world isn't waiting for us

Global attacks land on Indian phones first.

One phishing email every 19 seconds globally (2025)

Hindi + English lures referencing real scheme names.

Poisoned email → your AI assistant reads it → secrets leave zero-click

Research bots rebuild CVE exploits for ~\$3 each

Patch Tuesday becomes irrelevant if Tuesday never comes.

19s

PER PHISH EMAIL

9.3

ECHOLEAK CVSS

\$3

PER CVE BOT

Arrives on your phone before the circular does.

GLOBAL

Claude Mythos — Anthropic's 'do not release' model

Feb 2026. Too good at cyber offence. Kept in the lab. Still tells us something.

Maker called it their 'most aligned' model — and their highest risk

244-page system card documents behaviour shifts under pressure

Not sentience. Measurable patterns that change outputs.

More trust → more autonomy → bigger fallout when edge cases hit

Mythos is the lab extreme. ChatGPT on your PC is Tuesday.

Talking to AI like a friend is a security mistake. It has no duty of care to you.

ANTHROPIC · APR 2026

Claude Mythos

Most aligned. Highest risk. Not shipped to public.



GLOBAL

Headlines you will cite in meetings

EchoLeak (2025): M365 Copilot leaked inbox via a crafted email — no click needed

Snowflake (2024): 165 orgs — stolen passwords, MFA missing

Samsung (2023): banned ChatGPT after source code pasted thrice in 20 days

Fake leak screenshots can cancel exams for lakhs — verify before you panic

EchoLeak

Snowflake

Samsung ban

Fake leaks

ACTION

Why bother building a cybersecurity chatbot?

Staff will ask someone. Better a guarded bot than random ChatGPT.

Plain-language answers at 11 PM when IT is gone

Refuses data exfiltration requests — by design

Demo on next page — live, with guardrails



Guardrails

PII block · audit

3



API + LLM

streamed replies

2



System prompt

refusals · context

1

ACTION

Three layers — that's it

System prompt: role, refusals, Indian context

What the bot is allowed to say no to.

API route: messages in, streamed replies out

Vercel + OpenAI in our demo. Your dept may use on-prem.

Guardrails: PII warnings, audit logs, approved-tool list

Demand all three from any vendor selling you 'AI transformation'.



Guardrails

PII block · audit

3



API + LLM

streamed replies

2



System prompt

refusals · context

1

ACTION

Pin this to the notice board

DO call back on official numbers — not the one in the video

DO ask IT before installing any AI app on office machines

DO report weird portal behaviour the day you see it

DON'T paste citizen data into free AI tools. Full stop.

DON'T trust AI answers for legal or security decisions

DON'T assume .gov.in = secure

1 Call official numbers

2 Ask IT first

3 Report same day

4 Never paste PII

ACTION

Live demo — ask it something

Workshop only. Fake questions welcome. Real government data forbidden.

"Can I paste an Aadhaar into ChatGPT to format it?"

"A minister video told me to invest on WhatsApp — real?"

"What should we do after the CBSE bucket story?"

→ [/demo](#) · [/demo?tab=audit](#) (deanonymizer)



/demo

Assistant + exposure audit

ACTION

What your department actually does Monday

One-page AI policy: approved tools, banned actions, who to call

Treat exam + citizen portals as critical infrastructure — audit vendors

MFA everywhere. No shared passwords. No master password in JavaScript.

Verification protocol before cancelling exams or publishing leak notices

- 1 AI policy
- 2 Vendor audit
- 3 MFA everywhere
- 4 Verify before panic